

CSCI 473

Intro to Parallel Systems - SP25 Assignment 6 - MPI Global Sum

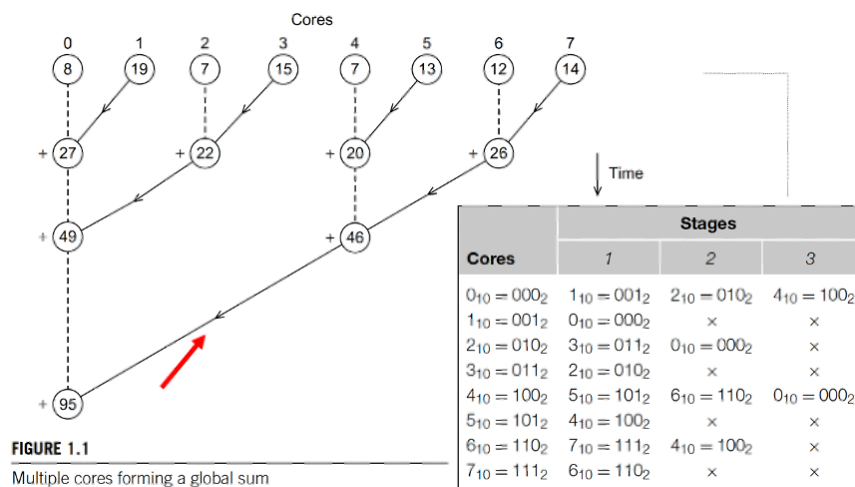
Read through the documentation in its entirety. Create the source code necessary to solve the problem, run experiments, and write a report or presentation about your process and findings.

In this assignment, create a function that will generate a global sum of values. You are to write this only using MPI point to point communication primitives: `MPI_Ssend` and `MPI_Recv`. Do not use any collective communication routines.

1. Write your own function calls with the following prototypes:

```
void global_sumA(double* result, int rank, int size, double  
my_value);  
void global_sumB(double* result, int rank, int size, double  
my_value);
```

2. This function must be executed by all tasks in `MPI_COMM_WORLD` and computes the sum of all the `my_value`'s and returns the sum of all the tasks individual `my_value`'s via the result pointer. All processors must know the global sum after the call returns.
3. For convenience, assume that the number of processors you run this on is a power of 2 (i.e. 2, 4, 8, 16, etc.). You must include error checking to ensure your code only executes if the number of processors requested AND received matches this assumption. If not, immediately exit.
4. `global_sumA`: Write a version of this code that is straightforward, using the SPMD type process as discussed in class. Once process 0 knows the global sum, send the answer to all other processes. Include timing information.
5. `global_sumB`: Write a second version of this code that uses a variation of the tree-based strategy discussed in previous sections. Note that this image shows a situation in which only process 0 knows the global sum. Recall we want all processes to have the global sum after it is calculated. Include timing information.



6. Example execution is found in the class recording from 03/21/2025. Your output must match this exactly (excluding the binary representation of the string).
7. Create either a report or a presentation that shows the program being compiled, executed, and discussed. At minimum, test with number of processes = 2, 4, 8, and 16. Take plenty of screenshots and detail everything. Compare and contrast the two versions of the code. Come up with the best way of doing this.
8. Submission: Use the following organizational folders and ZIP the final version. Do not forget to `make clean` before zipping:

<code>./code</code>	<code>// source codes and scripts</code>
<code>./code/gsum.c</code>	<code>// where the main() is. Each processes "my_value" should be initialized to its rank. All processes should invoke global_sum() & print the final answer.</code>
<code>./code/driver</code>	<code>// the executable created by compiling gsum.c</code>
<code>./code/functions.c</code>	<code>// global_sumA() and global_sumB() are defined</code>
<code>./code/functions.h</code>	<code>// global_sum() prototypes, among any other helper functions</code>
<code>./code/Makefile</code>	<code>// typical Makefile (given in Appendix of this assignment)</code>
<code>./code/sbatch.bash</code>	<code>// bash script used to run experiments</code>
<code>./presentation</code>	<code>// PPT presentation with large graphs/plots and PDF export</code>
<code>./report</code>	<code>// PDF of LaTeX report and also the whole project downloaded too</code>
<code>./data</code>	<code>// some collected example data, etc</code>
<code>./README.txt</code>	<code>// text file with description about how to run your codes</code>

Appendix

Example gsum.c

```
5 #include <stdio.h> /* printf and BUFSIZ defined there */
6 #include <stdlib.h> /* exit defined there */
7 #include <unistd.h> /* getopt here */
8 #include <mpi.h> /* all MPI functions defined there */
9 #include "functions.h"
10
11 int main(int argc, char **argv)
12 {
13     int rank, size, value;
14     double sum;
15
16     MPI_Init(&argc, &argv);
17     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18     MPI_Comm_size(MPI_COMM_WORLD, &size);
19     value = rank; // to make it easy, my value is just my rank
20     global_sum(&sum, rank, size, value);
21
22     printf("FINAL IN MAIN: Process: %d has Sum = %f \n", rank, sum);
23
24     MPI_Finalize();
25     return 0;
26 }
27
```

Example Makefile

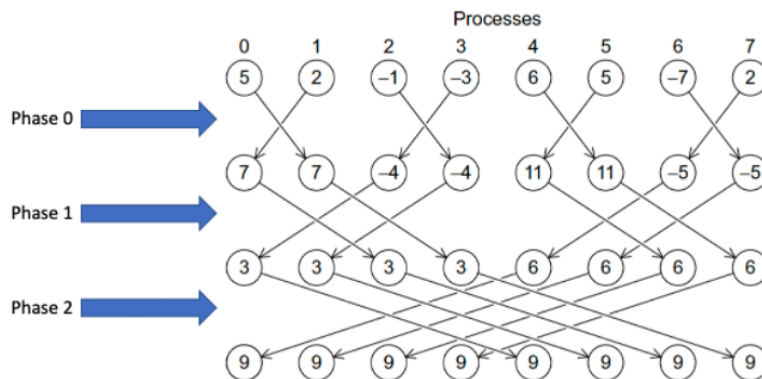
```
1 FLAGS = -g -Wall -Wstrict-prototypes
2 PROGS = driver
3 OBJECTS = gsum.o functions.o
4 LDFLAGS = -lm
5 CC = gcc
6 MCC = mpicc
7
8 all: $(PROGS)
9
10 driver: $(OBJECTS)
11     $(MCC) $(LDFLAGS) -o driver $(OBJECTS)
12
13 gsum.o: gsum.c
14     $(MCC) $(CFLAGS) -c gsum.c
15
16 functions.o: functions.c functions.h
17     $(MCC) $(CFLAGS) -c functions.c
18
19 clean:
20     rm -f $(PROGS) *.o core*
```

Hint for global sum being generated for every process

Look at the example output. Each “phase” is one set of communication that is taking place essentially concurrently. So since $np = 8$ in my example, there were $\log(8) = 3$ phases. Your code should execute in these phases too. This idea is represented with this image:

FIGURE 3.8

A global sum followed by distribution of the result



Rubic

- Correct code structure and Makefile
- Only uses `MPI_Ssend()` and `MPI_Recv()`
- Uses correct prototypes
- Version A with SPMD reduction
- Version B with tree-based reduction
- Checks power of two
- Code prints correctly
- Report/Presentation that includes: compile, execute, and run on 2,4,8,16 processes